

Pandas

January 11, 2021

1 Podstawy Analizy danych w Pythonie: pandas

1.1 10 stycznia 2021

Ostatnia część kursu Pythona będzie dotyczyć biblioteki **pandas**, która służy do analizy danych. Zaczniemy zatem od importu. Przeważnie bibliotekę skraca się do *pd*:

```
[65]: %matplotlib inline
import sys
import numpy as np
import matplotlib.pyplot as plt
```

Matplotlib is building the font cache; this may take a moment.

```
[168]: import pandas as pd
```

Pandas posiada dwie podstawowe struktury danych: * szereg (Series), * ramka danych (DataFrame).

Pandas pozwala na wczytanie danych z wielu formatów plików: * csv: `pd.read_csv` * json: `pd.read_json` * excel: `pd.read_excel` * SQL: `pd.read_sql`

Zobacz: https://pandas.pydata.org/pandas-docs/stable/reference/api/pandas.read_csv.html

```
[171]: df = pd.read_csv('/home/tomaszd/codes/big-data-python/labs01/gapminder.csv')
df.head(5)
```

```
[171]:
```

	Country	female_BMI	male_BMI	gdp	population	\
0	Afghanistan	21.07402	20.62058	1311.0	26528741.0	
1	Albania	25.65726	26.44657	8644.0	2968026.0	
2	Algeria	26.36841	24.59620	12314.0	34811059.0	
3	Angola	23.48431	22.25083	7103.0	19842251.0	
4	Antigua and Barbuda	27.50545	25.76602	25736.0	85350.0	

	under5mortality	life_expectancy	fertility
0	110.4	52.8	6.20
1	17.9	76.8	1.76
2	29.5	75.5	2.73
3	192.0	56.7	6.43

4	10.9	75.5	2.16
---	------	------	------

```
[11]: df = pd.read_excel('./bikes.xlsx', engine='openpyxl')
df.head()
```

```
[11]:
```

	start_date	start_station_code	end_date	\
0	2019-04-14 07:55:22	6001	2019-04-14 08:07:16	
1	2019-04-14 07:59:31	6411	2019-04-14 08:09:18	
2	2019-04-14 07:59:55	6097	2019-04-14 08:12:11	
3	2019-04-14 07:59:57	6310	2019-04-14 08:27:58	
4	2019-04-14 08:00:37	7029	2019-04-14 08:14:12	

	end_station_code	duration_sec	is_member
0	6132	713	1
1	6411	587	1
2	6036	736	1
3	6345	1680	1
4	6250	814	0

```
[31]: import sqlite3
select = "select * from logs"
with sqlite3.connect('astro-timeline.sqlite3') as db:
    astro_timeline = pd.read_sql(select, db, parse_dates=['datetime'])
astro_timeline.head()
```

```
[31]:
```

	id	datetime	level	\
0	1	1969-07-14 21:00:00+00:00	INFO	
1	2	1969-07-16 13:31:53+00:00	WARNING	
2	3	1969-07-16 13:33:23+00:00	DEBUG	
3	4	1969-07-16 13:34:44+00:00	WARNING	
4	5	1969-07-16 13:35:17+00:00	DEBUG	

	message
0	Terminal countdown started
1	S-IC engine ignition (#5)
2	Maximum dynamic pressure (735.17 lb/ft^2)
3	S-II ignition
4	Launch escape tower jettisoned

```
[26]: df = pd.read_csv('/home/tomaszd/codes/big-data-python/labs01/gapminder.csv',
↳ index_col='Country')
df.head()
```

```
[26]:
```

	female_BMI	male_BMI	gdp	population	\
Country					
Afghanistan	21.07402	20.62058	1311.0	26528741.0	
Albania	25.65726	26.44657	8644.0	2968026.0	
Algeria	26.36841	24.59620	12314.0	34811059.0	
Angola	23.48431	22.25083	7103.0	19842251.0	
Antigua and Barbuda	27.50545	25.76602	25736.0	85350.0	

	under5mortality	life_expectancy	fertility
Country			
Afghanistan	110.4	52.8	6.20
Albania	17.9	76.8	1.76
Algeria	29.5	75.5	2.73
Angola	192.0	56.7	6.43
Antigua and Barbuda	10.9	75.5	2.16

```
[28]: df = pd.read_csv("./titanic_train.tsv", sep='\t', index_col='PassengerId')
df.head()
```

```
[28]:
```

	Survived	Pclass	\
PassengerId			
1	0	3	
2	1	1	
3	1	3	
4	1	1	
5	0	3	

	Name	Sex	Age	\
PassengerId				
1	Braund\t Mr. Owen Harris	male	22.0	
2	Cumings\t Mrs. John Bradley (Florence Briggs T...	female	38.0	
3	Heikkinen\t Miss. Laina	female	26.0	
4	Futrelle\t Mrs. Jacques Heath (Lily May Peel)	female	35.0	
5	Allen\t Mr. William Henry	male	35.0	

	SibSp	Parch	Ticket	Fare	Cabin	Embarked
PassengerId						
1	1	0	A/5 21171	7.2500	NaN	S
2	1	0	PC 17599	71.2833	C85	C
3	0	0	STON/O2. 3101282	7.9250	NaN	S
4	1	0	113803	53.1000	C123	S
5	0	0	373450	8.0500	NaN	S

```
[173]: import sqlite3
select = "select * from logs"
with sqlite3.connect('astro-timeline.sqlite3') as db:
```

```

astro_timeline = pd.read_sql(select, db, parse_dates=['datetime'])

astro_timeline.to_csv('astro-timeline.csv')
# astro_timeline.head()

```

```

[173]:
   id      datetime      level \
0    1 1969-07-14 21:00:00+00:00  INFO
1    2 1969-07-16 13:31:53+00:00 WARNING
2    3 1969-07-16 13:33:23+00:00  DEBUG
3    4 1969-07-16 13:34:44+00:00 WARNING
4    5 1969-07-16 13:35:17+00:00  DEBUG

      message
0  Terminal countdown started
1      S-IC engine ignition (#5)
2  Maximum dynamic pressure (735.17 lb/ft^2)
3      S-II ignition
4  Launch escape tower jettisoned

```

2 Szeregi (pd.Series)

```

[175]: dane = pd.Series([5, 6, 7, 8])
print(dane)

```

```

0    5
1    6
2    7
3    8
dtype: int64

```

```

[36]: members = {'April': 211819, 'May': 682758, 'June': 737011, 'July': 779511,
                 'August': 673790, 'September': 673790, 'October': 444177, 'November':
                 ↪ 136791,
                 }
dane = pd.Series(list(members.values()))
print(dane)

```

```

0    211819
1    682758
2    737011
3    779511
4    673790
5    673790
6    444177
7    136791
dtype: int64

```

Czym różni się szereg od listy? Szereg danych posiada indeks, czyli klucz, dzięki któremu możemy zindyfikować dane. Domyślnie, indeks jest ciągiem liczb zaczynających się od zera. Nie musi tak być, możemy podczas tworzenia przekazać również indeks:

```
[176]: dane = pd.Series(list(members.values()), index=members.keys())
print(dane)
```

```
April      211819
May        682758
June       737011
July       779511
August     673790
September  673790
October    444177
dtype: int64
```

```
[177]: dane = pd.Series(list(members.values()), index=members.keys())
print(dane['April'])
```

```
211819
```

```
[179]: dane = pd.Series(list(members.values()), index=members.keys())
keys = ['April', 'September']
dane[keys]
```

```
[179]: April      211819
September  673790
dtype: int64
```

```
[182]: dane = pd.Series(list(members.values()), index=members.keys())
print(dane['June': 'September'])
```

```
June      737011
July      779511
August    673790
September 673790
dtype: int64
```

```
[184]: dane = pd.Series(list(members.values()), index=members.keys())

dane['June'] = 333000
dane
```

```
[184]: April      211819
May        682758
June       333000
July       779511
```

```
August      673790
September   673790
October     444177
dtype: int64
```

```
[53]: print(len(dane))
      print(dane.shape)
```

```
8
(8,)
```

Przeważnie zbiory danych, na których pracujemy są duże. Stąd, próba ich wyświetlenia może okazać się karkołomna lub nawet niemożliwa. Czasem chcemy tylko zobaczyć pogląd. Do tego służą dwie metody: **head** i **tail**, które zwrócą nam kilka pierwszych lub ostatnich wierszy z szeregu:

```
[185]: dane.head()
```

```
[185]: April      211819
      May        682758
      June       333000
      July       779511
      August     673790
      dtype: int64
```

```
[186]: dane.tail()
```

```
[186]: June       333000
      July       779511
      August     673790
      September   673790
      October     444177
      dtype: int64
```

```
[188]: dane.sample(5)
```

```
[188]: June       333000
      July       779511
      September   673790
      May        682758
      April      211819
      dtype: int64
```

Szeregi są dostosowane do analizy danych. Np. udostępniają prosty sposób do uzyskania podstawowych statystyk:

```
[189]: dane = pd.Series([1, 3, 2, 3, 1, 1, 2, 3, 2, 3])
      print("Średnia:", dane.mean())
      print("Mediana:", dane.median())
```

Średnia: 2.1

Mediana: 2.0

Jak i inne przydatne funkcje:

```
[191]: dane = pd.Series([1, 3, 2, 3, 1, 1, 2, 3, 2, 3])  
  
# print("Zbiór wartości:", dane.unique())  
print(dane.value_counts())
```

```
3    4
```

```
2    3
```

```
1    3
```

```
dtype: int64
```

Metoda `value_counts` zwraca nam szereg danych, który możemy wykorzystać do dalszych badań. Na przykład, żeby wyświetlić 5 najczęściej występujących wartości, możemy napisać:

```
[192]: print(dane.value_counts().head(1))
```

```
3    4
```

```
dtype: int64
```

Żeby uzyskać wszystkie podstawowe statystyki, możemy wywołać metodę `describe`:

```
[62]: print(dane.describe())
```

```
count    10.000000
```

```
mean      2.100000
```

```
std       0.875595
```

```
min       1.000000
```

```
25%       1.250000
```

```
50%       2.000000
```

```
75%       3.000000
```


```
max       3.000000
```

```
dtype: float64
```

A żeby wyświetlić je w postaci wykresu:

```
[66]: dane.hist()
```

```
[66]: <AxesSubplot:>
```



```
[194]: dane = pd.Series(list(members.values()), index=members.keys())
```

```
dane.index
```

```
[194]: Index(['April', 'May', 'June', 'July', 'August', 'September', 'October'],
dtype='object')
```

```
[196]: dane = pd.Series(list(members.values()), index=members.keys())
```

```
dane.values
```

```
[196]: array([211819, 682758, 737011, 779511, 673790, 673790, 444177])
```

Jeżeli chcemy zmienić cały szereg przy pomocy funkcji, możemy wykorzystać metodę `map`:

```
[198]: def cube(x):
        return x / 1000

print(dane.map(cube))
# dane
```

April	211.819
May	682.758
June	737.011
July	779.511
August	673.790


```
September    673.790
October       444.177
dtype: float64
```

Uwaga: w Pythonie istnieją funkcje lambda, które można tu wykorzystać.

```
[200]: members = {'April': 211819, 'May': 682758, 'June': 737011, 'July': 779511,
                  'August': 673790, 'September': 673790, 'October': 444177, 'November':
                  ↪ 136791, 'December': None
                }

dane = pd.Series(list(members.values()), index=members.keys())
print(dane.tail())
```

```
August        673790.0
September     673790.0
October       444177.0
November      136791.0
December      NaN
dtype: float64
```

```
[203]: dane.notna()
```

```
[203]: April          True
      May            True
      June           True
      July           True
      August         True
      September      True
      October        True
      November       True
      December       False
dtype: bool
```

```
[205]: members = {'April': 211819, 'May': 682758, 'June': 737011, 'July': None,
                  'August': 673790, 'September': None, 'October': 444177, 'November':
                  ↪ 136791, 'December': None
                }

dane = pd.Series(list(members.values()), index=members.keys())
print(dane.head())
```

```
April        211819.0
May          682758.0
June         737011.0
July         NaN
August       673790.0
dtype: float64
```

```
[206]: dane.dropna()
```

```
[206]: April      211819.0
      May       682758.0
      June      737011.0
      August    673790.0
      October   444177.0
      November  136791.0
      dtype: float64
```

```
[208]: dane.fillna(dane.mean())
```

```
[208]: April      211819.000000
      May       682758.000000
      June      737011.000000
      July      481057.666667
      August    673790.000000
      September 481057.666667
      October   444177.000000
      November  136791.000000
      December  481057.666667
      dtype: float64
```

```
[88]: dane.interpolate()
```

```
[88]: April      211819.0
      May       682758.0
      June      737011.0
      July      705400.5
      August    673790.0
      September 558983.5
      October   444177.0
      November  136791.0
      December  136791.0
      dtype: float64
```

```
[209]: members = {'April': 211819, 'May': 682758, 'June': 737011, 'July': 779511,
      ↪            'August': 673790,
      'September': 673790, 'October': 444177, 'November': 136791}

members_series = pd.Series(list(members.values()), index=members.keys())

print(members_series + 1000)
```

```
April      212819
May        683758
```

```

June          738011
July           780511
August         674790
September      674790
October        445177
November       137791
dtype: int64

```

```

[92]: members = {'April': 211819, 'May': 682758, 'June': 737011, 'July': 779511,
↳ 'August': 673790,
        'September': 673790, 'October': 444177, 'November': 136791}

occasionals = {'April': 32058, 'May': 147898, 'June': 171494, 'July': 194316,
↳ 'August': 206809,
        'September': 140492, 'October': 53596, 'November': 10516}

members_series = pd.Series(list(members.values()), index=members.keys())
occasionals_series = pd.Series(list(occasionals.values()), index=occasionals.
↳ keys())

all_series = members_series + occasionals_series
all_series

```

```

[92]: April          243877
May           830656
June          908505
July          973827
August        880599
September     814282
October       497773
November      147307
dtype: int64

```

```

[94]: members = {'April': 211819, 'May': 682758, 'June': 737011, 'July': 779511,
↳ 'August': 673790,
        'September': 673790, 'October': 444177}

occasionals = {'May': 147898, 'June': 171494, 'July': 194316, 'August': 206809,
↳ 'September': 140492, 'October': 53596, 'November': 10516}

members_series = pd.Series(list(members.values()), index=members.keys())
occasionals_series = pd.Series(list(occasionals.values()), index=occasionals.
↳ keys())

all_series = members_series + occasionals_series
all_series

```

```
[94]: April          NaN
      August      880599.0
      July        973827.0
      June        908505.0
      May         830656.0
      November     NaN
      October     497773.0
      September   814282.0
      dtype: float64
```

```
[210]: members = {'April': 211819, 'June': 737011, 'July': 779511, 'August': 673790,
                  'September': 673790, 'October': 444177}

occasionals = {'May': 147898, 'June': 171494, 'July': 194316, 'August': 206809,
               'September': 140492, 'October': 53596, 'November': 10516}

members_series = pd.Series(list(members.values()), index=members.keys())
occasionals_series = pd.Series(list(occasionals.values()), index=occasionals.
                               ↪keys())

all_series = members_series.add(occasionals_series, fill_value=0)
all_series
```

```
[210]: April          211819.0
      August      880599.0
      July        973827.0
      June        908505.0
      May         147898.0
      November     10516.0
      October     497773.0
      September   814282.0
      dtype: float64
```

2.1 Ramki danych (pd.DataFrame)

```
[211]: members = {'April': 211819, 'May': 682758, 'June': 737011, 'July': 779511,
                  ↪'August': 673790,
                  'September': 673790, 'October': 444177}

occasionals = {'May': 147898, 'June': 171494, 'July': 194316, 'August': 206809,
               'September': 140492, 'October': 53596, 'November': 10516}

df = pd.DataFrame({'members': members, 'occasionals': occasionals})

df.head()
```

```
[211]:
```

	members	occasionals
April	211819.0	NaN
May	682758.0	147898.0
June	737011.0	171494.0
July	779511.0	194316.0
August	673790.0	206809.0

```
[212]: data = [{'members': 682758, 'occasionals': 147898},
               {'members': 737011, 'occasionals': 171494},
               {'members': 779511, 'occasionals': 194316}]

df = pd.DataFrame(data)

df.head()
```

```
[212]:
```

	members	occasionals
0	682758	147898
1	737011	171494
2	779511	194316

```
[213]: data = [(682758, 147898), (737011, 171494), (779511, 194316)]

df = pd.DataFrame(data)

df.head()
```

```
[213]:
```

	0	1
0	682758	147898
1	737011	171494
2	779511	194316

```
[215]: data = [(682758, 147898), (737011, 171494), (779511, 194316)]

df = pd.DataFrame(data)
df.columns = ['members', 'occasionals']

df.head()
```

```
[215]:
```

	members	occasionals
0	682758	147898
1	737011	171494
2	779511	194316

```
[216]: data = [(682758, 147898), (737011, 171494), (779511, 194316)]

df = pd.DataFrame(data, columns=['members', 'occasionals'])
```

```
df.head()
```

```
[216]:   members  occasionals
0    682758      147898
1    737011      171494
2    779511      194316
```

```
[114]: members = {'April': 211819, 'May': 682758, 'June': 737011, 'July': 779511,
    ↪ 'August': 673790,
        'September': 673790, 'October': 444177}

occasionals = {'May': 147898, 'June': 171494, 'July': 194316, 'August': 206809,
    ↪ 'September': 140492, 'October': 53596, 'November': 10516}

df = pd.DataFrame({'members': members, 'occasionals': occasionals})

df.index
```

```
[114]: Index(['April', 'May', 'June', 'July', 'August', 'September', 'October',
    ↪ 'November'],
    ↪ dtype='object')
```

```
[219]: members = {'April': 211819, 'May': 682758, 'June': 737011, 'July': 779511,
    ↪ 'August': 673790,
        'September': 673790, 'October': 444177}

occasionals = {'May': 147898, 'June': 171494, 'July': 194316, 'August': 206809,
    ↪ 'September': 140492, 'October': 53596, 'November': 10516}

df = pd.DataFrame({'members': members, 'occasionals': occasionals})

df.values
```

```
[219]: array([[211819.,    nan],
    ↪ [682758., 147898.],
    ↪ [737011., 171494.],
    ↪ [779511., 194316.],
    ↪ [673790., 206809.],
    ↪ [673790., 140492.],
    ↪ [444177.,  53596.],
    ↪ [    nan, 10516.]])
```

```
[220]: df.head()
```

```
[220]:   members  occasionals
April    211819.0         NaN
May      682758.0    147898.0
```

June	737011.0	171494.0
July	779511.0	194316.0
August	673790.0	206809.0

```
[221]: df.tail()
```

```
[221]:
```

	members	occasionals
July	779511.0	194316.0
August	673790.0	206809.0
September	673790.0	140492.0
October	444177.0	53596.0
November	NaN	10516.0

```
[222]: print(df.max())
```

```
members      779511.0
occasionals   206809.0
dtype: float64
```

```
[119]: print(df.describe())
```

	members	occasionals
count	7.000000	7.000000
mean	600408.000000	132160.142857
std	201535.422382	73336.753852
min	211819.000000	10516.000000
25%	558983.500000	97044.000000
50%	673790.000000	147898.000000
75%	709884.500000	182905.000000
max	779511.000000	206809.000000

```
[120]: print(df.sample(3))
```

	members	occasionals
June	737011.0	171494.0
November	NaN	10516.0
July	779511.0	194316.0

```
[224]: # df.members
```

```
df['members']
```

```
[224]:
```

April	211819.0
May	682758.0
June	737011.0
July	779511.0
August	673790.0
September	673790.0

```
October      444177.0
November      NaN
Name: members, dtype: float64
```

```
[125]: df[['members', 'occasionals']]
```

```
[125]:      members  occasionals
April      211819.0          NaN
May         682758.0      147898.0
June        737011.0      171494.0
July        779511.0      194316.0
August       673790.0      206809.0
September   673790.0      140492.0
October     444177.0       53596.0
November      NaN       10516.0
```

```
[225]: df.loc['August']
```

```
[225]: members      673790.0
occasionals    206809.0
Name: August, dtype: float64
```

```
[128]: df.loc[['August', 'September']]
```

```
[128]:      members  occasionals
August      673790.0      206809.0
September   673790.0      140492.0
```

```
[129]: df.loc['June': 'September']
```

```
[129]:      members  occasionals
June        737011.0      171494.0
July         779511.0      194316.0
August       673790.0      206809.0
September    673790.0      140492.0
```

```
[130]: df.loc['June': 'September', 'members']
```

```
[130]: June      737011.0
July         779511.0
August        673790.0
September     673790.0
Name: members, dtype: float64
```

```
[226]: df.at['June', 'members']
```

```
[226]: 737011.0
```



```
[133]: df.at['June', 'members'] = 123456
df.at['June', 'members']
```

```
[133]: 123456.0
```

```
[228]: df['all_rides'] = df['members'] + df['occasionals']

print(df.head())
```

	members	occasionals	all_rides
April	211819.0	NaN	NaN
May	682758.0	147898.0	830656.0
June	737011.0	171494.0	908505.0
July	779511.0	194316.0	973827.0
August	673790.0	206809.0	880599.0

```
[229]: df.loc['December'] = [0, 0, 0]

print(df.tail())
```

	members	occasionals	all_rides
August	673790.0	206809.0	880599.0
September	673790.0	140492.0	814282.0
October	444177.0	53596.0	497773.0
November	NaN	10516.0	NaN
December	0.0	0.0	0.0

```
[151]: df.drop('April')
```

```
[151]:
```

	members	occasionals	all_rides
May	682758.0	147898.0	830656.0
June	737011.0	171494.0	908505.0
July	779511.0	194316.0	973827.0
August	673790.0	206809.0	880599.0
September	673790.0	140492.0	814282.0
October	444177.0	53596.0	497773.0
November	NaN	10516.0	NaN
December	0.0	0.0	0.0

```
[154]: df.drop('members', axis='columns') # df.drop(columns='members')
```

```
[154]:
```

	occasionals	all_rides
April	NaN	NaN
May	147898.0	830656.0
June	171494.0	908505.0
July	194316.0	973827.0
August	206809.0	880599.0

September	140492.0	814282.0
October	53596.0	497773.0
November	10516.0	NaN
December	0.0	0.0

```
[155]: df.transpose()
```

```
[155]:
```

	April	May	June	July	August	September \
members	211819.0	682758.0	737011.0	779511.0	673790.0	673790.0
occasionals	NaN	147898.0	171494.0	194316.0	206809.0	140492.0
all_rides	NaN	830656.0	908505.0	973827.0	880599.0	814282.0

	October	November	December
members	444177.0	NaN	0.0
occasionals	53596.0	10516.0	0.0
all_rides	497773.0	NaN	0.0

```
[231]: df = pd.read_csv("./titanic_train.tsv", sep='\t', index_col='PassengerId')
df.head()
```

```
[231]:
```

	Survived	Pclass \
PassengerId		
1	0	3
2	1	1
3	1	3
4	1	1
5	0	3

	Name	Sex	Age \
PassengerId			
1	Braund\t Mr. Owen Harris	male	22.0
2	Cumings\t Mrs. John Bradley (Florence Briggs T...	female	38.0
3	Heikkinen\t Miss. Laina	female	26.0
4	Futrelle\t Mrs. Jacques Heath (Lily May Peel)	female	35.0
5	Allen\t Mr. William Henry	male	35.0

	SibSp	Parch	Ticket	Fare	Cabin	Embarked
PassengerId						
1	1	0	A/5 21171	7.2500	NaN	S
2	1	0	PC 17599	71.2833	C85	C
3	0	0	STON/O2. 3101282	7.9250	NaN	S
4	1	0	113803	53.1000	C123	S
5	0	0	373450	8.0500	NaN	S

```
[18]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 891 entries, 1 to 891
Data columns (total 11 columns):
Survived      891 non-null int64
Pclass        891 non-null int64
Name          891 non-null object
Sex           891 non-null object
Age           714 non-null float64
SibSp         891 non-null int64
Parch         891 non-null int64
Ticket        891 non-null object
Fare          891 non-null float64
Cabin         204 non-null object
Embarked      889 non-null object
dtypes: float64(2), int64(4), object(5)
memory usage: 83.5+ KB
```

```
[22]: print(df.shape)
```

```
(891, 11)
```

```
[232]: # df['Survived'] == 1
df.loc[df['Survived'] == 1].head()
```

```
[232]: PassengerId
1      False
2       True
3       True
4       True
5      False
...
887    False
888     True
889    False
890     True
891    False
Name: Survived, Length: 891, dtype: bool
```

```
[140]: df['Survived'] == 1
```

```
[140]: PassengerId
1      False
2       True
3       True
4       True
5      False
...
887    False
```

```

888     True
889     False
890     True
891     False
Name: Survived, Length: 891, dtype: bool

```

```

[144]: survived = df['Survived'] == 1
first_class = df['Pclass'] == 1

# df[survived & first_class]
# df[survived | first_class]
df[~first_class]

```

```

[144]:
      Survived  Pclass      Name \
PassengerId
1           0        3  Braund\t Mr. Owen Harris
3           1        3  Heikkinen\t Miss. Laina
5           0        3  Allen\t Mr. William Henry
6           0        3  Moran\t Mr. James
8           0        3  Palsson\t Master. Gosta Leonard
...         ...      ...
885          0        3  Sutehall\t Mr. Henry Jr
886          0        3  Rice\t Mrs. William (Margaret Norton)
887          0        2  Montvila\t Rev. Juozas
889          0        3  Johnston\t Miss. Catherine Helen "Carrie"
891          0        3  Dooley\t Mr. Patrick

```

```

      Sex  Age  SibSp  Parch      Ticket    Fare Cabin \
PassengerId
1    male  22.0    1     0  A/5 21171    7.2500   NaN
3  female  26.0    0     0  STON/O2. 3101282    7.9250   NaN
5    male  35.0    0     0    373450    8.0500   NaN
6    male   NaN    0     0    330877    8.4583   NaN
8    male   2.0    3     1    349909   21.0750   NaN
...     ...   ...   ...   ...
885   male  25.0    0     0  SOTON/OQ 392076    7.0500   NaN
886  female  39.0    0     5    382652   29.1250   NaN
887   male  27.0    0     0    211536   13.0000   NaN
889  female   NaN    1     2  W./C. 6607   23.4500   NaN
891   male  32.0    0     0    370376    7.7500   NaN

```

```

      Embarked
PassengerId
1           S
3           S
5           S
6           Q

```

```

8          S
...
885        S
886        Q
887        S
889        S
891        Q

```

[675 rows x 11 columns]

```
[237]: df[['Sex', 'Survived']].groupby('Sex').mean()
```

```

[237]:      Survived
Sex
female  0.742038
male    0.188908

```

```
[166]: df.groupby(['Sex', 'Pclass']).mean()
```

```

[166]:      Survived      Age   SibSp  Parch    Fare
Sex  Pclass
female 1      0.968085  34.611765  0.553191  0.457447  106.125798
      2      0.921053  28.722973  0.486842  0.605263   21.970121
      3      0.500000  21.750000  0.895833  0.798611   16.118810
male   1      0.368852  41.281386  0.311475  0.278689   67.226127
      2      0.157407  30.740707  0.342593  0.222222   19.741782
      3      0.135447  26.507589  0.498559  0.224784   12.661633

```

```
[ ]:
```